ASGE · AMBIVALENCE-CENTRED SEMANTIC GEOMETRY ENGINE

Integration Guide

Quickstart, schema, error handling, and operational notes for engineers integrating ASGE into an analytical pipeline.

AUDIENCE	VERSION	ISSUED
Engineering	**1.0**	**May 2026**

01 · OVERVIEW

One endpoint. A complete analysis.

ASGE exposes one operation: submit a passage, receive the full three-stage analysis of ambivalence structures across the timeline of the text. The response is a single JSON object that carries lexical pair density, peak concentration moments, the dominant semantic-conflict theme, a structured analytical summary, and a graduated alert layer.

There is no SDK to install, no message queue to provision, no callback URL to register. A single HTTPS request returns the full analysis in seconds.

02 · AUTHENTICATION

Session tokens.

API access is provisioned per organisation against a tier (Professional, API / Integration, Enterprise). Sandbox access is issued separately on request for evaluation. Write to `growth.business@llulllab.com` with the tier you intend, the contracting entity, and the expected use case. A session token is issued by return.

The session token is passed in the `X-Session-Token` header on every request. Tokens are scoped to a single organisation and may be rotated on request. Token lifecycle, including expiry and rotation policy, is set at access provisioning.

> **CUSTODY** Treat the session token as a production secret. Storage in a managed secret vault is recommended. Rotation on personnel change is recommended. Compromise should be reported to `growth.business@llulllab.com` for immediate revocation.

03 · THE REQUEST

Submit a passage for analysis.

The endpoint is `POST https://api.llulllab.com/asge/v1/analyse`. Content type is `application/json`. The minimum viable request carries the text, a language code, and the session token.

```
# Submit a passage for ambivalence-structure analysis
curl -X POST https://api.llulllab.com/asge/v1/analyse \
  -H "X-Session-Token: $ASGE_TOKEN" \
  -H "Content-Type: application/json" \
  -d '{
    "text":     "The text to analyse. Paste a transcript, an interview, a speech, or any natura
    "language": "en",
    "mode":     "text",
    "context":  {"source": "interview-2026-04-12"}
  }'
```

REQUEST FIELDS

FIELD	TYPE	DESCRIPTION
text	string · required	The passage to analyse. Up to 10,000 words per submission. Line breaks segment temporal units.
language	enum · required	One of `en`, `de`, `fr`, `es`, `tr`, `ja`, `he`, `fa`.
mode	enum · optional	`text` (default) or `microphone`. The latter is reserved for sessions originating from a live recording context.
context	object · optional	Optional metadata kept alongside the case record (source, author, date, location). Up to 4 KB.

TIP Line breaks segment temporal units. For passages that lack natural minute structure (essays, articles), break the text into paragraphs roughly equivalent to one minute of speech and the temporal density map will read naturally.

04 · THE RESPONSE

A structured analysis, returned in one call.

A successful response carries the full analysis as a single JSON object. The most useful fields, in order of how analysts typically consume them: the peak minute, the dominant theme, the conflict summary, the alert level, the trend timeline.

```
{
  "success": true,
  "analysis": {
    "pairs": [
      { "minute": 3, "type": "classic",    "context": "..." },
      { "minute": 7, "type": "contextual", "context": "..." }
    ],
    "trend": [
      { "minute": 1, "level": "low",    "note": "..." },
      // minute-by-minute
    ],
    "peak_minute":      4,
    "dominant_theme":   "freedom ↔ constraint",
    "conflict_summary": "Two to three sentences naming the sustained tensions.",
    "alert_level":      3
  },
  "stats": {
    "word_count":    1820,
    "pair_count":    24,
    "classic_count": 17,
    "context_count": 7,
    "pairs_per_min": 2.4
  }
}
```

RESPONSE FIELDS

FIELD	TYPE	DESCRIPTION
success	boolean	True when the three-stage pipeline completed.
analysis.pairs	array	Each entry: minute, type (`classic` or `contextual`), and a short context fragment.
analysis.trend	array	Minute-by-minute density with level and a one-line analytical note.
analysis.peak_minute	integer	The minute index where ambivalence density is highest.
analysis.dominant_theme	string	The dominant semantic-conflict theme, expressed as an opposition.
analysis.conflict_summary	string	Two to three sentences naming the sustained tensions.
analysis.alert_level	integer	Graduated alert from 1 to 4. Higher values correspond to denser concentrations.
stats	object	Counts and densities: words, pairs, classic and contextual splits, pairs per minute.

05 · STATUS CODES

Six outcomes worth handling.

● 200 · success ● 400 · validation ● 401 · unauthenticated ● 413 · payload_too_large
● 429 · quota_exceeded ● 503 · transient_failure

VALIDATION (400)

The request was rejected before analysis. Common causes: text below the minimum length, invalid language enum, malformed JSON. Not billed.

UNAUTHENTICATED (401)

The session token is missing, malformed, or revoked. Verify the `X-Session-Token` header. Not billed.

PAYLOAD TOO LARGE (413)

The text exceeds 10,000 words. Split the passage and submit in batches; consider whether segmenting changes the analytical question.

QUOTA EXCEEDED (429)

The tier's monthly included allowance, or a short-window rate limit, was exhausted. Overage pricing applies on tiers that permit it. Not billed when rejected.

TRANSIENT FAILURE (503)

An upstream component briefly degraded. Safe to retry once after a short backoff. Auto-retry is performed internally; what reaches the client is a definitive verdict. Not billed.

---

06 · OPERATIONAL NOTES

What you can rely on.

TRANSPORT

HTTPS only, encrypted in transit. Connection keep-alive is honoured.

LATENCY

Typical end-to-end latency for a completed analysis depends on text length. Short passages (under 1,000 words) typically complete in under 20 seconds; longer passages scale with the number of temporal segments produced. The analysis is not designed for synchronous user-facing flows.

AUDIT TRAIL

Every analysis is persisted with full lineage: input payload, detected pairs, temporal trend, theme, summary, and alert level. Retention and retrieval terms are set per tier at access provisioning.

DATA RESIDENCY

Production infrastructure is hosted in the European Union. Deployment options for the Enterprise tier are discussed at contracting.

ASGE · Integration Guide · v1.0 · Llull Lab · Studio Llull · Paris · growth.business@llulllab.com · llulllab.com/ASGE

This document is published as supporting material for the ASGE service. The infrastructure underneath is proprietary and is not part of the public contract.